Generating Minimal Models
for
Geometric Theories

A Major Qualifying Project Report
submitted to the Faculty of

Worcester Polytechnic Institute

in partial fulfillment of
the requirements for the degree of
Bachelor of Science

by

....

Michael Ficarra

on

11^{th} October, 2010

....

Daniel Dougherty
professor, project advisor

Abstract

This paper describes a method, referred to as the chase, for generating jointly minimal models for a geometric theory. A minimal model for a theory is a model for which there exists a homomorphism to any other model that can satisfy the theory. These models are useful in solutions to problems in many practical applications, including but not limited to firewall configuration examination, protocol analysis, and access control evaluation. Also described is a Haskell implementation of the chase and its development process and design decisions.

i

Table of Contents

1   Introduction

This document details the

1.1   Goals

The two main goals of this Major Qualifying Project are:

1. first, to implement an algorithm known as "the chase" accurately and with a well-defined, usable interface

2. second, to use the chase implementation for a real-world application: generating models used in analysis of a specific protocol

Secondary goals include implementing various optimizations and integrating the chase implementation into a program that can take advantage of the functionality it provides.

1.2   The Chase

2 Technical Background

2.1 Models

A *model* $\mathbb{M}$ is a construct that consists of:

- a set, referenced as $|\mathbb{M}|$, called the *universe* or *domain* of $\mathbb{M}$

- a set of pairings of a *predicate* and a non-negative integral arity

- for each predicate R with arity k, a relation $R_k^{\mathbb{M}} \subseteq |\mathbb{M}|$

It is important to distinguish the predicate, which is just a symbol, from the relation that it refers to when paired with its arity. The relation itself is a set of tuples of members from the universe.

2.2 First-order Logic

First-order logic, also called *predicate logic*, is a formal logic system. A first-order logic formula is defined inductively by

- if R is a relation symbol of arity k and each of $x_0 \ldots x_{k-1} \in \vec{x}$ is a variable, then $R[\vec{x}]$ is a formula, specifically an *atomic formula*

- if x and y are variables, then $x = y$ is a formula

- $\top$ and $\bot$ are formulæ

- if α is a formula, then $(\neg\alpha)$ is a formula

- if α and β are formulæ, then $(\alpha \wedge \beta)$ is a formula

- if α and β are formulæ, then $(\alpha \vee \beta)$ is a formula

- if α and β are formulæ, then $(\alpha \to \beta)$ is a formula

- if α is a formula and x is a variable, then $(\forall x : \alpha)$ is a formula

- if α is a formula and x is a variable, then $(\exists x : \alpha)$ is a formula

For our purposes, this logic system will not contain any constant symbols or function symbols, which are commonly included in first-order logic. A function $f(x, y)$ can be encoded as the geometric logic formula $F[x, y, z1] \wedge F[x, y, z2]-> z1 = z2$. Constant symbols can be encoded as relations with an arity of zero, for example $C[]$.

A shorthand notation may sometimes be used which omits either the left or right side of an implication and denotes $(\top \to \sigma)$ and $(\sigma \to \bot)$ respectively. If α is a formula and $\vec{x}$ is a set of variables of size k, then $(\forall \vec{x} : \alpha)$ is $(\forall x_0 \ldots \forall x_{k-1} : \alpha)$. If α is a formula and $\vec{x}$ is a set of variables of size k, then $(\exists \vec{x} : \alpha)$ is $(\exists x_0 \ldots \exists x_{k-1} : \alpha)$.

2.3 Variable Binding

The set of free variables in a formula is defined inductively as follows

- any variable occuring in an atomic formula is a free variable

- the set of free variables in $\top$ and $\bot$ is $\emptyset$

- the set of free variables in $x = y$ is $\{x, y\}$

- the set of free variables in $\neg \alpha$ is $free(\alpha)$

- the set of free variables in $\alpha \wedge \beta$ is $free(\alpha) \cup free(\beta)$

- the set of free variables in $\alpha \vee \beta$ is $free(\alpha) \cup free(\beta)$

- the set of free variables in $\alpha \to \beta$ is $free(\alpha) \cup free(\beta)$

- the set of free variables in $\forall x : \alpha$ is $free(\alpha) - \{x\}$

- the set of free variables in $\exists x : \alpha$ is $free(\alpha) - \{x\}$

A formula α is a *sentence* if $free(\alpha) = \emptyset$.

2.4 Environment

An *environment* λ for a model $\mathbb{M}$ is a function from a variable v to a domain member e where $e \in |\mathbb{M}|$. The syntax $\lambda_{[v \mapsto a]}$ denotes the environment $\lambda'(x)$ that returns a when $x = v$ and returns $\lambda(x)$ otherwise.

2.5 Satisfiability

A model $\mathbb{M}$ is said to satisfy a formula σ in an environment λ, denoted $\mathbb{M} \models_\lambda \sigma$ and read "under λ, σ is true in $\mathbb{M}$", when

- σ is a relation symbol R and $R[\lambda(a_0), \ldots, \lambda(a_n)] \in \mathbb{M}$ where a is a set of variables

- σ is of the form $\neg \alpha$ and $\mathbb{M} \nvDash_\lambda \alpha$

- σ is of the form $\alpha \wedge \beta$ and both $\mathbb{M} \models_\lambda \alpha$ and $\mathbb{M} \models_\lambda \beta$

- σ is of the form $\alpha \vee \beta$ and either $\mathbb{M} \models_\lambda \alpha$ or $\mathbb{M} \models_\lambda \beta$

- σ is of the form $\alpha \rightarrow \beta$ and either $\mathbb{M} \not\models_\lambda \alpha$ or $\mathbb{M} \models_\lambda \beta$

- σ is of the form $\forall x : \alpha$ and for every $x' \in |\mathbb{M}|$, $\mathbb{M} \models_{\lambda[x \mapsto x']} \alpha$

- σ is of the form $\exists x : \alpha$ and for at least one $x' \in |\mathbb{M}|$, $\mathbb{M} \models_{\lambda[x \mapsto x']} \alpha$

The notation $\mathbb{M} \models \sigma$ (no environment specification) means that, under the empty environment l, $\mathbb{M} \models_l \sigma$.

A model $\mathbb{M}$ satisfies a set of formulæ Σ under an environment λ if for every σ such that $\sigma \in \Sigma$, $\mathbb{M} \models_\lambda \sigma$. This is denoted as $\mathbb{M} \models_\lambda \Sigma$ and read "$\mathbb{M}$ is a model of Σ".

2.6 Entailment

Given an environment λ, a set of formulæ Σ is said to *entail* a formula σ ($\Sigma \models_\lambda \sigma$) if the set of all models satisfied by Σ under λ is a subset of the set of all models satisfying σ under λ.

The notation used for satisfiability and entailment is very similar, in that the operator used ($\models$) is the same, but they can be distinguished by the type of left operand.

2.7 Homomorphisms

A *homomorphism* from $\mathbb{A}$ to $\mathbb{B}$ is a function $h : |\mathbb{A}| \rightarrow |\mathbb{B}|$ such that, for each relation symbol R and tuple $\langle a_0, \ldots, a_n \rangle$ where $a \subseteq |\mathbb{A}|$, $\langle a_0, \ldots, a_n \rangle \in R^{\mathbb{A}}$ implies $\langle h(a_0), \ldots, h(a_n) \rangle \in R^{\mathbb{B}}$.

A homomorphism h is also a *strong homomorphism* if, for each relation symbol R and tuple $\langle a_0, \ldots, a_n \rangle$ where $a \subseteq |\mathbb{A}|$, $\langle a_0, \ldots, a_n \rangle \in R^{\mathbb{A}}$ if and only if $\langle h(a_0), \ldots, h(a_n) \rangle \in R^{\mathbb{B}}$.

The notation $\mathbb{M} \preceq \mathbb{N}$ means that there exists a homomorphism $h : \mathbb{M} \rightarrow \mathbb{N}$. The identity function is a homomorphism from any model $\mathbb{M}$ to itself. Homomorphisms have the property that $\mathbb{A} \preceq \mathbb{B} \wedge \mathbb{B} \preceq \mathbb{C}$ implies $\mathbb{A} \preceq \mathbb{C}$.

However, $\mathbb{M} \preceq \mathbb{N} \wedge \mathbb{N} \preceq \mathbb{M}$ does not imply that $\mathbb{M} = \mathbb{N}$, but instead that $\mathbb{M}$ and $\mathbb{N}$ are *homomorphically equivalent*. For example, fix two models $\mathbb{M}$ and $\mathbb{N}$ that are equivalent except that $\mathbb{N}$ has one more domain member than $\mathbb{M}$. Both $\mathbb{M} \preceq \mathbb{N}$ and $\mathbb{N} \preceq \mathbb{M}$ are true, yet $\mathbb{M} \neq \mathbb{N}$. Homomorphic Equivalence between a model $\mathbb{M}$ and a model $\mathbb{N}$ is denoted $\mathbb{M} \simeq \mathbb{N}$.

Given models $\mathbb{M}$ and $\mathbb{N}$ where $\mathbb{M} \preceq \mathbb{N}$ and a formula in positive-existential form σ, if $\mathbb{M} \models \sigma$ then $\mathbb{N} \models \sigma$.

A homomorphism $h : \mathbb{A} \rightarrow \mathbb{B}$ is also an *isomorphism* when h is 1:1 and onto and the inverse function $h^{-1} : \mathbb{B} \rightarrow \mathbb{A}$ is a homomorphism.

2.8   Minimal Models

Minimal models, also called *universal* models, are models for a theory with the special property that there exists a homomorphism from the minimal model to any other model that satisfies the theory. Intuitively, minimal models have no unnecessary entities or relations and thus display the least amount of constraint necessary to satisfy the theory for which they are minimal.

More than one minimal model may exist for a given theory, and not every theory must have a minimal model. **give examples**.

A set of models $\mathcal{M}$ is said to be *jointly minimal* for a set of formulæ Σ when every model $\mathbb{N}$ such that $\mathbb{N} \models \Sigma$ has a homomorphism from a model $\mathbb{M} \in \mathcal{M}$ to $\mathbb{N}$.

2.9   Positive Existential Form

Formulæ in *positive existential form* are constructed using only conjunctions ($\wedge$), disjunctions ($\vee$), existential quantifications ($\exists$), tautologies ($\top$), contradictions ($\bot$), equalities, and relations.

Negation of a relation R with arity k can be implemented by introducing another relation R' with arity k, adding two formulæ of the form $R \wedge R' \rightarrow \bot$ and $\top \rightarrow R \vee R'$, and using R' where $\neg R$ would be used.

2.10   Geometric Logic

Geometric logic formulæ are implicitly universally quantified implications between positive existential formulæ. More specifically, a geometric logic formula is of the form

$$\forall\, (free(F_L) \cup free(F_R)) : F_L \rightarrow F_R$$

where $free$ is the function that returns the set of all free variables for a given formula and both F_L and F_R are are first-order logic formulæ in positive existential form.

A set of geometric logic formulæ is called a *geometric theory*.

It is convention to treat a positive existential formula σ as $\top \rightarrow \sigma$ when expecting a geometric logic formula. It is also convention to treat a negated positive existential formula $\neg \sigma$ as $\sigma \rightarrow \bot$.

Examples of geometric logic formulæ:

reflexivity	$\top \rightarrow R[x, x]$
symmetry	$R[x, y] \rightarrow R[y, x]$
transitivity	$R[x, y] \wedge R[y, z] \rightarrow R[x, z]$

3 The Chase

The *chase* is a function that, when given a gemoetric theory, will generate a set of jointly minimal models for that theory. More specifically, if $\mathcal{U}$ is the set of all models obtained from an execution of the chase over a geometric theory T, for any model $\mathbb{M}$ such that $\mathbb{M} \models T$, there is a homomorphism from some model $\mathbb{U} \in \mathcal{U}$ to $\mathbb{M}$.

Geometric logic formulæ are used by the chase because they have the useful property where adding any relations or domain members to a model that satisfies a geometric logic formula will never cause the formula to no longer be satisified. This is particularly helpful when trying to create a model that satisifies all formulæ in a geometric theory.

There are three types of runs of the chase:

- a non-empty result in finite time

- an empty result in finite time

- an infinite run, with possible return dependent on implementation

3.1 Algorithm

The chase algorithm begins by defining an empty list of models to return $\mathcal{D}$ and a list of pending models $\mathcal{P}$ containing a model with an empty domain and an empty set of facts. A single input is provided to the chase: the geometric theory Σ.

Before anything else, the chase algorithm sorts Σ by the number of disjunctions on the right side of each formula's implication. Formulæ with zero disjunctions will be ordered first, followed by those with one disjunction, and finally those with more than one disjunction. Within each sorting classification, formulæ should remain in their original order.

For every model $\mathbb{M} \in \mathcal{P}$, the chase loops through every formula $\sigma \in \Sigma$. For any σ and environment λ such that $\mathbb{M} \not\models_\lambda \sigma$, $\mathbb{M}$ has domain members and facts added to it **as described in ¡algorithm below¿**. If $\mathbb{M} \models \Sigma$ instead, remove $\mathbb{M}$ from $\mathcal{P}$ and add it to $\mathcal{D}$.

algorithm for how to make a model satisfy a formula

Finally, the chase algorithm returns $\mathcal{D}$, a list of jointly minimal models for its input theory Σ.

3.2 Examples

Define Σ as the following geometric theory.

$$\top \quad \to \quad \exists\, x, y : R[x, y] \tag{1}$$
$$R[x, y] \quad \to \quad (\exists\, z : Q[x, z]) \vee P \tag{2}$$
$$Q[x, y] \quad \to \quad (\exists\, z : R[x, z]) \vee (\exists\, z : R[z, y]) \tag{3}$$
$$P \quad \to \quad \bot \tag{4}$$
$$\tag{5}$$

The following three chase runs show the different types of results depending on which disjunct the algorithm attempts to satisfy when a disjunction is encountered.

1. A non-empty result in finite time:

$$\emptyset \; \mapsto \; \{ \quad a, b \quad | \quad R[a, b] \quad \}$$
$$\mapsto \; \{ \quad a, b, c \quad | \quad R[a, b], Q[a, c] \quad \}$$

Since the left side of **(1)** is always satisfied, but its right side is not, domain members a and b and fact $R[a, b]$ are added to the initially empty model to satisfy **(1)**. The left side of **(2)** holds, but the right side does not, so one of the disjuncts $\exists\, z : Q[x, z]$ or $P[x]$ is chosen to be satisfied. Assuming the left operand is chosen, x will already have been assigned to a and a new domain member c and a new fact $Q[a, c]$ will be added to satisfy **(2)**. With the current model, all rules hold under any environment. Therefore, this model is in the set of jointly minimal models for our theory.

2. An empty result in finite time:

$$\emptyset \; \mapsto \; \{ \quad a, b \quad | \quad R[a, b] \quad \}$$
$$\mapsto \; \{ \quad a, b, c \quad | \quad R[a, b], P[a, c] \quad \}$$
$$\mapsto \; \{ \quad a, b, c \quad | \quad R[a, b], P[a, c], \bot \quad \}$$
$$\mapsto \; \varepsilon$$

Again, domain members a and c and fact $R[a, b]$ are added to the initial model to satisfy **(1)**. This time, when attempting to satisfy **(2)**, the right side is chosen and P is added to the set of facts. After adding this new fact, rule **(4)** no longer holds; its left side is satisfied, but its right side does not hold for all of the bindings for which it is satisfied. When we attempt to satisfy the right side of **(4)**, it is found to be a contradiction and therefore unsatisfiable. Since this model can never satisfy this theory, the chase fails.

3. An infinite run:

$$\emptyset \; \mapsto \; \{ \quad a, b \quad | \quad R[a, b] \quad \}$$
$$\mapsto \; \{ \quad a, b, c \quad | \quad R[a, b], Q[a, c] \quad \}$$
$$\mapsto \; \{ \quad a, \ldots, d \quad | \quad R[a, b], Q[a, c], R[d, c] \quad \}$$
$$\mapsto \; \{ \quad a, \ldots, e \quad | \quad R[a, b], Q[a, c], R[d, c], Q[d, e] \quad \}$$
$$\mapsto \; \{ \quad a, \ldots, f \quad | \quad R[a, b], Q[a, c], R[d, c], Q[d, e], R[f, e] \quad \}$$
$$\mapsto \; \ldots$$

Like in the example above that returned a non-empty, finite result, the first two steps add domain members a, b, and c and facts $R[a, b]$ and $Q[a, c]$. The left side

of the implication in **(3)** now holds, but the right side does not. In order to make the right side hold, one of the disjuncts needs to be satisfied. If the right disjunct is chosen, a new domain member d and a new relation $R[d, c]$ will be added. This will cause the left side of the implication in **(2)** to hold for $R[d, c]$, but the right side will no hold for the same binding. $Q[d, e]$ will be added, and this loop will continue indefinitely unless a different disjunct is chosen in **(2)** or **(3)**.

4   Haskell Chase Implementation

The goal of the implementation of the chase is to deterministically find all possible outcomes of the chase. It does this by forking and taking all paths when encountering a disjunct rather than nondeterministically choosing one disjunct to satisfy.

The results from the attempts to satisfy each disjunct are returned as a list. The returned list will not contain an entry for runs that return no model, and will merge lists returned from runs that themselves encountered a disjunct. The lazy evaluation of Haskell allows a user to access members of the returned list even though some chase runs have not returned a value.

Appendix B contains the chase-running portions of the implementation.

4.1   Operation

The first step of the chase implementation is to make sure that each formula of the given theory can be represented as a geometric logic formula. If a formula φ can not be coerced to a geometric logic formula, the chase tries to coerce it into one by applying following rules recursively:

- $\neg \alpha \wedge \neg \beta \mapsto \alpha \vee \beta$

- $\neg \alpha \vee \neg \beta \mapsto \alpha \wedge \beta$

- $\neg \neg \alpha \mapsto \alpha$

- $\neg \alpha \rightarrow \beta \mapsto \alpha \vee \beta$

- $\alpha \rightarrow \neg \beta \mapsto \alpha \wedge \beta$

- $\neg (\forall \, \vec{x} : \neg \alpha) \mapsto \exists \vec{x} : \alpha$

All other constructs are preserved. If this transformed formula is still not in positive existential form, an error is thrown.

The chase function then sorts the input formulæ by the number of disjunctions on the right side of the implication. it allows branches to terminate without growing to an unnecessarily (and possibly infinitely) large size. This step will cause each branch of the algorithm to finish in less time, as they are likely to halt before branching yet again. The formulæ are not sorted purely by absolute number of disjunctions on the right side, but by whether there are zero, one, or many disjunctions. This is done to avoid unnecessary re-ordering for no gain because formulæ with no disjunctions or only a single disjunct on the right side are more likely to cause a branch to stop growing than one with many

disjunctions. Likewise, formulæ with zero disjunctions are more likely to cause a branch to halt than those with one or more disjunctions.

Once the input formulæ are sorted, the *chase* function begins processing a *pending* list, which is initially populated with a single model that has an empty domain and no facts.

For each *pending* model, each formula is evaluated to see if it holds in the model for all environments. If an envionment is found that does not satisfy the model, the model and environment in which the formula did not hold is passed to the *satisfy* function, along with the formula that needs to be satisfied. The list of models returned from *satisfy* is merged into the *pending* list, and the result of running *chase* on the new *pending* list is returned. If, however, the model holds for all formulæ in the theory and all possible associated environments, it is concatenated with the result of running the chase on the rest of the models in the *pending* list.

The *satisfy* function performs a pattern match on the type of formula given. Assuming *satisfy* is given a model $\mathbb{M}$, an environment λ, and a formula φ, *satisfy* will behave as

outlined in the following algorithm.

Algorithm: satisfy :: Model → Environment → Formula → [Model]

return **switch** φ **do**

> **case** $\top$ return a list containing $\mathbb{M}$
>
> **case** $\bot$ return an empty list
>
> **case** $x = y$ return a list containing $quotient(\mathbb{M})$
>
> **case** $\alpha \vee \beta$ return $satisfy(\alpha) \cup satisfy(\beta)$
>
> **case** $\alpha \wedge \beta$
>
>> create an empty list r
>>
>> **foreach** $model\ m\ in\ satisfy(\alpha)$ **do**
>>
>>> union r with $satisfy(\beta)$
>>
>> return r
>
> **case** $\alpha \rightarrow \beta$
>
>> **if** $\mathbb{M}_\lambda \models \alpha$ **then** return $satisfy(\beta)$
>>
>> **else** return an empty list
>
> **case** $R[\vec{x}]$
>
>> define a new model $\mathbb{N}$ where $|\mathbb{N}| = |\mathbb{M}|$
>>
>> add a new member ω to $|\mathbb{N}|$
>>
>> **forall the** $P_\mathbb{M}$ **do** $P_\mathbb{N} = P_\mathbb{M}$
>>
>> define $R_\mathbb{N}[x_0 \ldots x_n]$ as $R_\mathbb{M}[\lambda(x_0) \ldots \lambda(x_n)]$
>>
>> **foreach** $v \in \vec{x}$ **do**
>>
>>> **if** $v \notin \lambda$ **then** λ becomes $\lambda_{v \mapsto \omega}$
>>
>> return a list containing $\mathbb{N}$
>
> **case** $\exists\ \vec{x} : \alpha$
>
>> **if** $\vec{x} = \emptyset$ **then** recurse on α
>>
>> **if** $|\mathbb{M}| \neq \emptyset\ and\ \exists\ v' \in |\mathbb{M}| : (\lambda' = \lambda_{x_0 \mapsto v'}\ and\ \mathbb{M} \models_{\lambda'} \alpha)$ **then**
>>
>>> return a list containing $\mathbb{M}$
>>
>> **else**
>>
>>> define a new model $\mathbb{N}$ where $|\mathbb{N}| = |\mathbb{M}|$
>>>
>>> add a member ω to $|\mathbb{N}|$ such that $\omega \notin |\mathbb{N}|$
>>>
>>> **forall the** $R_\mathbb{M}$ **do** $R_\mathbb{N} = R_\mathbb{M}$
>>>
>>> define $\kappa = \lambda_{x_0 \mapsto \omega}$
>>>
>>> using model $\mathbb{N}$ and environment κ, return $satisfy(\exists\ \{x_1 \ldots x_n\} : \alpha)$

4.2   Input Format

Input to the parser must be in a form parsable by the following context-free grammar. Terminals are denoted by a `monospace style` and nonterminals are denoted by an *obliquestyle*. The greek letter ε matches a zero-length list of tokens. Patterns that match non-literal terminals are defined in the table following the grammar.

$$
\begin{array}{lll}
program & : \varepsilon \\
& |\ exprList\ optNEWLINE \\[1em]
exprList & :\ expr \\
& |\ exprList\ \texttt{NEWLINE}\ expr \\[1em]
expr & :\ \texttt{TAUTOLOGY} \\
& |\ \texttt{CONTRADICTION} \\
& |\ expr\ \texttt{OR}\ expr \\
& |\ expr\ \texttt{AND}\ expr \\
& |\ \texttt{NOT}\ expr \\
& |\ expr\ \texttt{->}\ expr \\
& |\ \texttt{->}\ expr \\
& |\ atomic \\
& |\ \texttt{VARIABLE EQ VARIABLE} \\
& |\ \texttt{FOR\_ALL}\ argList\ optCOLON\ expr \\
& |\ \texttt{THERE\_EXISTS}\ argList\ optCOLON\ expr \\
& |\ \texttt{(}\ expr\ \texttt{)} \\
& |\ \texttt{[}\ expr\ \texttt{]} \\[1em]
atomic & :\ \texttt{PREDICATE}\ index \\[1em]
index & :\ \texttt{(}\ argList\ \texttt{)} \\
& |\ \texttt{[}\ argList\ \texttt{]} \\[1em]
argList & :\ arg \\
& |\ argList\ \texttt{,}\ arg \\[1em]
arg & :\ \texttt{VARIABLE} \\[1em]
optCOLON & :\ \varepsilon \\
& |\ \texttt{:} \\[1em]
optNEWLINE & :\ \varepsilon \\
& |\ \texttt{NEWLINE}
\end{array}
$$

Input Pattern	Terminal
`|`	`OR`
`&`	`AND`
`!`	`NOT`
`=`	`EQ`
`[Tt]autology`	`TAUTOLOGY`
`[Cc]ontradiction`	`CONTRADICTION`
`[\r\n]+`	`NEWLINE`
`[a-z][A-Za-z0-9_']*`	`VARIABLE`
`[A-Z][A-Za-z0-9_']*`	`PREDICATE`
`For[Aa]ll`	`FOR_ALL`
`Exists`	`THERE_EXISTS`

Comments are removed at the lexical analysis step and have no effect on the input to the parser. Single-line comments begin with either a hash (`#`) or double-dash (`--`). Multiline comments begin with `/*` and are terminated by `*/`.

4.3   Options

4.3.1   I/O

When no options are given to the executable output by Haskell, it expects input from stdin and outputs models in a human-readable format to stdout. To take input from a file instead, pass the executable the `-i` or `--input` option followed by the filename.

To output models to numbered files in a directory, pass the `-o` or `--output` option along with an optional directory name. The given directory does not have to exist. If the output directory is omitted, it defaults to "`./models`".

Using the `-o` or `--output` options will change the selection for output format to a machine-readable format. To switch ouput formats at any time, pass the `-h` or `-m` flags for human-readable and machine-readable formats respectively.

4.3.2   Tracing

not yet implemented

4.4   Future Considerations

This section details areas of possible improvement/development that were unable to be explored during the timeframe allotted for the project.

4.4.1 Better Data Structures

This project was started with no knowledge of the most desirable implementation language, Haskell. Because of this, some less-than-optimal data structures were used to hold data that should really be in a Data.Map or Data.Set. One such example of this is with the truth table contained in relations. This truth table should probably be implemented as a Data.Map. Also, instead of Domains being a list of DomainMember, it would probably be better if a Domain was a Data.Set.

4.4.2 Broader use of the Maybe Monad

4.4.3 Advanced Rewriters / Simplifiers

talk about purple book and a reverse pullquants

5 An Extended Application: Cryptographic Protocol Analysis

The chase can be used for protocol analysis. A common technique for the analysis of protocols involves identifying the *essentially different* runs of the protocol. These essentially different protocol runs are analagous to minimal models. When a protocol is described using geometric logic, the chase can find such minimal models.

The protocol can then be analysed for characteristics such as the existence of security violations or other unexpected behaviour.

5.1 Background

5.1.1 Strand Spaces

The strand space formalism was developed as a method of formally reasoning about cryptographic protocols. Participants in a protocol run are represented by *strands*. A single physical entity can be respresented as multiple strands if they play many roles in the protocol. A strand is made up of a sequence of nodes where every node either sends or receives a message.

Adversary strands represent the actions of a participant that attempts to use the protocol for a purpose other than the one for which it was originally intended. Adversary strands are not bound by the rules defined by the protocol; they manipulate messages being sent and received by non-adversarial strands.

The capabilities of the adversary strands are given by the Dolev-Yao Threat Model **references**. Adversary strands are able to perform precisely five operations:

pairing the pairing of two terms

unpairing the extraction of a term from a pair

encryption given a key k and a plaintext m, the construction of the ciphertext $\{|m|\}_k$ by encrypting m with k

decryption given a ciphertext $\{|m|\}_k$ and its decryption key k^{-1}, the extraction of the plaintext m

generation the generation of an original term, which is not assumed to be secure

Pairing, encryption, and generation are construction operations. Decryption and unpairing are deconstruction operations.

5.1.2 Cremers' Algorithm

Cremers' Algorithm adds constraints to adversarial actions to allow one to infer the possible shapes of protocol runs.

Cremers gives two major constraints. A protocol is *efficient* if an adversary always takes a message from the earliest point at which it appears. A protocol is *normal* when the adversary always performs zero or more deconstruction operations followed by zero or more construction operations, except in the case of constructing decryption keys.

Cremers proved that these constraints do not limit the capabilties of an adversary **reference**.

An important insight used in Cremers' algorithm, called *chaining*, states that terms in messages received from an adversary strand always originate in a non-adversarial strand.

5.2 The Problem

Protocol reasearchers want to be able to programmatically reason about cryptographic protocols. A common technique for this is to find a set of essentially different classes of protocol runs. This can be accomplished by finding minimal models of a geometric logic representation of the protocol. This happens to be precisely the problem the chase solves.

5.3 The Solution: Minimal Models

Given a theory $\mathcal{T}$, the jointly minimal models $\mathcal{M}$ that the chase outputs are representative of all models because there always exists a homomorphism from some model $\mathbb{M} \in \mathcal{M}$ to any model that satisfies $\mathcal{T}$. Each model that satisfies $\mathcal{T}$ represents a class of runs of the protocol. The set of all models output by the chase represents every possible run of the protocol. Finding every possible run of the protocol is prohibitive because there are infinitely many. Because the set of models is countably infinite, they can be enumerated, but can not be listed.

5.4 Designing An Analagous Theory

In order to create a geometric theory describing a protocol, the formulæ that define strand spaces, normilisation, efficiency, and chaining must be derived. The formulæ defining the protocol must be combined with this scaffolding to create a theory that can be used to infer the possible runs of the protocol.

The *half-duplex protocol* was chosen to be used as an example. This protocol involves two participants, Alice and Bob. The protocol specifies that the following actions take place:

1. Alice sends Bob a nonce that she generated, encrypted with Bob's public key

2. Bob receives the encrypted nonce

3. Bob replies to Alice with the decypted nonce

4. Alice receives Bob's message

The geometric logic rules that model this protocol were generated manually by direct translation into geometric logic. Ideally, the process of generating geometric logic formulæ from protocols should be done automatically.

5.5   The Results

The chase was run on the logic representation of the half-duplex protocol. A single model was returned during the execution of the algorithm, which was manually stopped before natural completion. This model, like all models returned by the chase, satisfies the input theory, and belongs to a set of jointly minimal models for the theory. The returned model denotes a run of the protocol which contains no advesary strands and is a correct execution of the protocol.

A   Table of Syntax

syntax	definition
f^{-1}	the inverse function of f
$R[a_0, a_1, a_2]$	a *relation* of: relation symbol R, arity 3, and tuple $\langle a_0, a_1, a_2 \rangle$
$\top$	a *tautological formula*; one that will always hold
$\bot$	a *contradictory formula*; one that will never hold
$\rho = \tau$	given assumed environment λ, $\lambda(\rho) = \lambda(\tau)$
$\neg\alpha$	α does not hold
$\alpha \wedge \beta$	both α and β hold
$\alpha \vee \beta$	either α or β hold
$\alpha \rightarrow \beta$	either α does not hold or β holds
$\forall x : \alpha$	for each member of the domain as x, α holds
$\forall \vec{x} : \alpha$	**FIXME:** for each $x_i \in x$, $\forall x_i : \alpha$ holds
$\exists x : \alpha$	for at least one member of the domain as x, α holds
$\exists \vec{x} : \alpha$	**FIXME:** for each $x_i \in x$, $\exists x_i : \alpha$ holds
$\lambda[x \mapsto y]$	the environment λ with variable x mapped to domain member y
$\mathbb{M} \models_l \sigma$	$\mathbb{M}$ *is a model of* σ under environment l
$\mathbb{M} \models \sigma$	$\mathbb{M} \models_l \sigma$ given any environment l
$\mathbb{M} \models_l \Sigma$	for each $\sigma \in \Sigma$, $\mathbb{M} \models_l \sigma$
$\mathbb{M} \models \Sigma$	for each $\sigma \in \Sigma$, $\mathbb{M} \models \sigma$
$\Sigma \models \sigma$	Σ *entails* σ
$\mathbb{M} \preceq \mathbb{N}$	there exists a *homomorphism* $h : |\mathbb{M}| \rightarrow |\mathbb{N}|$
$\mathbb{M} \simeq \mathbb{N}$	$\mathbb{M}$ and $\mathbb{N}$ are *homomorphically equivalent*

B   Chase code

TODO: make the long lines short so they fit

```
 1  module Chase where
 2  import Parser
 3  import Helpers
 4  import qualified Debug.Trace
 5  import Data.List
 6
 7  -- trace x = id
 8  trace = Debug.Trace.trace
 9
10  verify :: Formula -> Formula
11  -- verifies that a formula is in positive existential form and performs some
12  -- normalization on implied/constant implications
13  verify formula = case formula of
14      Implication a b -> Implication (pef a) (pef b)
15      Not f -> Implication (pef f) Contradiction
16      _ -> Implication Tautology (pef formula)
17
18  order :: [Formula] -> [Formula]
19  --
20  order formulae = sortBy (\a b ->
21      let extractRHS = (\f -> case f of; (Implication lhs rhs) -> rhs; _ -> f) in
22      let (rhsA,rhsB) = (extractRHS a, extractRHS b) in
23      let (lenA,lenB) = (numDisjuncts rhsA, numDisjuncts rhsB) in
24      let (vA,vB) = (length (variables a), length (variables b)) in
25      if lenA == lenB || lenA > 1 && lenB > 1 then
26          if vA == vB then EQ
27          else
28              if vA < vB then LT
29              else GT
30      else
31          if lenA < lenB then LT
32          else GT
33      ) formulae
34
35  chase :: [Formula] -> [Model]
36  -- a wrapper for the chase' function to hide the model identity and theory
37  -- manipulation
38  chase theory = nub $ chase' (order $ map verify theory) [([],[])]
39
40  chase' :: [Formula] -> [Model] -> [Model]
41  -- runs the chase algorithm on a given theory, manipulating the given list of
42  -- models, and returning a list of models that satisfy the theory
43  chase' _ [] = trace "  but it is impossible to make the model satisfy the theory by adding to it" [
44  chase' theory pending = concatMap (branch theory) pending
45
46  branch :: [Formula] -> Model -> [Model]
47  --
48  branch theory model =
49      let reBranch = chase' theory in
50      trace ("running chase on " ++ show model) $
51      case findFirstFailure model theory of
52          Just newModels ->
53              trace ("  at least one formula does not hold for model " ++ showModel model) $
54              reBranch newModels
55          Nothing -> -- represents no failures
56              trace ("all formulae in theory hold for current model") $
57              trace ("returning model " ++ showModel model) $
58              [model]
59
60  findFirstFailure :: Model -> [Formula] -> Maybe [Model]
```

20

```haskell
61   ---
62   findFirstFailure model [] = Nothing --- no failure found
63   findFirstFailure model@(domain,relations) (f:ormulae) =
64       let self = findFirstFailure model in
65       let bindings = allBindings (freeVariables f) domain [] in
66       trace ("__checking_formula:_(v:" ++ show (length$variables f) ++ ")_(fv:" ++ show (length$freeVa
67       if holds model (UniversalQuantifier (freeVariables f) f) then self ormulae
68       else Just $ findFirstBindingFailure model f bindings
69
70   findFirstBindingFailure :: Model -> Formula -> [Environment] -> [Model]
71   ---
72   findFirstBindingFailure _ (Implication a Contradiction) _ = []
73   findFirstBindingFailure model formula@(Implication a b) (e:es) =
74       let self = findFirstBindingFailure model formula in
75       if holds' model e formula then self es
76       else
77           trace ("__attempting_to_satisfy_(" ++ show b ++ ")_with_env_" ++ show e) $
78           satisfy model e b
79
80   satisfy :: Model -> Environment -> Formula -> [Model]
81   ---
82   satisfy model env formula =
83       let (domain,relations) = model in
84       let domainSize = length domain in
85       let self = satisfy model in
86       case formula of
87           Tautology -> [model]
88           Contradiction -> []
89           Or a b -> union (self env a) (self env b)
90           And a b -> concatMap (\m -> satisfy m env b) (self env a)
91           Equality v1 v2 -> case (lookup v1 env, lookup v2 env) of
92               (Just v1, Just v2) -> [quotient model v1 v2]
93               _ -> error("Could_not_look_up_one_of_\"" ++ show v2 ++ "\"_or_\"" ++ show v2 ++ "\"_in_env
94           Atomic predicate vars ->
95               let newRelationArgs = genNewRelationArgs env vars (fromIntegral (length domain)) in
96               let newRelation = mkRelation predicate (length vars) [newRelationArgs] in
97               let newModel = mkModel (mkDomain domainSize) (mergeRelation newRelation relations) in
98               trace ("____adding_new_relation:_" ++ show newRelation) $
99               [newModel]
100          ExistentialQuantifier [] f -> self env f
101          ExistentialQuantifier (v:vs) f ->
102              let f' = ExistentialQuantifier vs f in
103              let nextDomainMember = fromIntegral $ (length domain) + 1 in
104              if (domain /= []) && (any (\d -> holds' model (hashSet env v d) f') domain) then
105                  trace ("____" ++ show formula ++ "_already_holds") $
106                  [model]
107              else
108                  trace ("____adding_new_domain_element_" ++ show nextDomainMember ++ "_for_variable_" ++
109                  satisfy (mkDomain nextDomainMember,relations) (hashSet env v nextDomainMember) f'
110          _ -> error ("formula_not_in_positive_existential_form:_" ++ show formula)
111
112  genNewRelationArgs :: Environment -> [Variable] -> DomainMember -> [DomainMember]
113  --- for each Variable in the given list of Variables, retrieves the value
114  --- assigned to it in the given environment, or the next domain element if it
115  --- does not exist
116  genNewRelationArgs env [] domainSize = []
117  genNewRelationArgs env (v:vs) domainSize =
118      let self = genNewRelationArgs env vs in
119      case lookup v env of
120          Just v' -> v' : self domainSize
121          _ -> domainSize+1 : self (domainSize+1)
```

21

References

[1] A Cottrell, *Word Processors: Stupid and Inefficient,*

         `www.ecn.wfu.edu/~cottrell/wp.html`